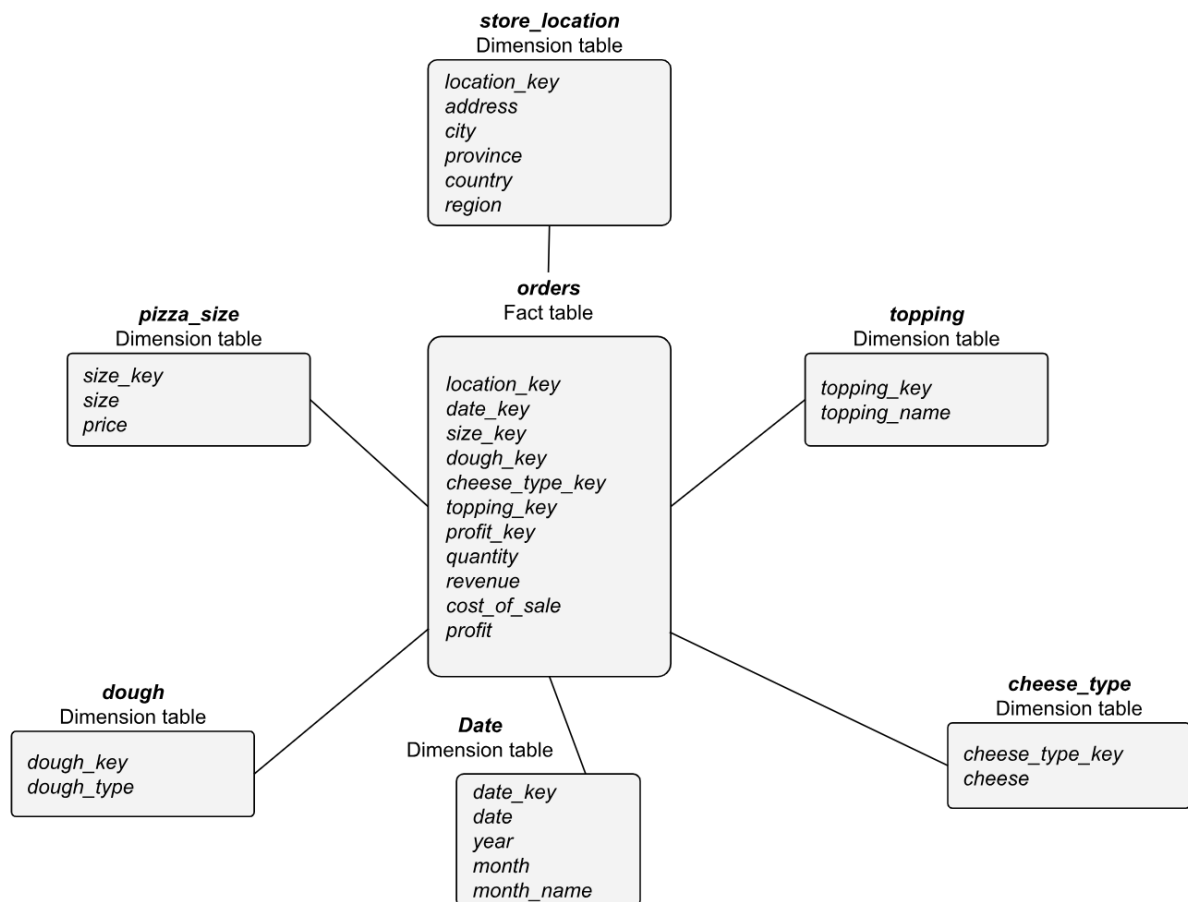


**Data Warehousing, OLAP Operations, Slicing and Dicing Activities****Part A****1a. Star schema**

The following star schema includes Orders (Fact) as well as all of its dimensions. A date dimension has been added as well for the ease of calculations in subsequent questions.

It should be noted that the size dimension (labeled *pizza\_size*) also includes the prices of each pizza size. For the purpose of this activity, we are assuming that the price of the pizza is dependent on its size. Therefore, the other attributes of the pizza (topping, dough, cheese) do not contribute to the total price of the pizza. Revenue and profit attributes are added to the schema as they calculate the revenue and profit of the order, using another attribute called cost-of-sale (i.e.,  $\text{profit} = \text{revenue} - \text{cost of sale}$ ).



## 1b. Snowflake schema

In the above star schema the store\_location dimension table is not normalized. Hence, we have introduced the city dimension table and country dimension table to make a snowflake, to normalize the snowflake schema and reduce redundancies. Such a table makes it easy to maintain and saves processing capabilities and thus resources.

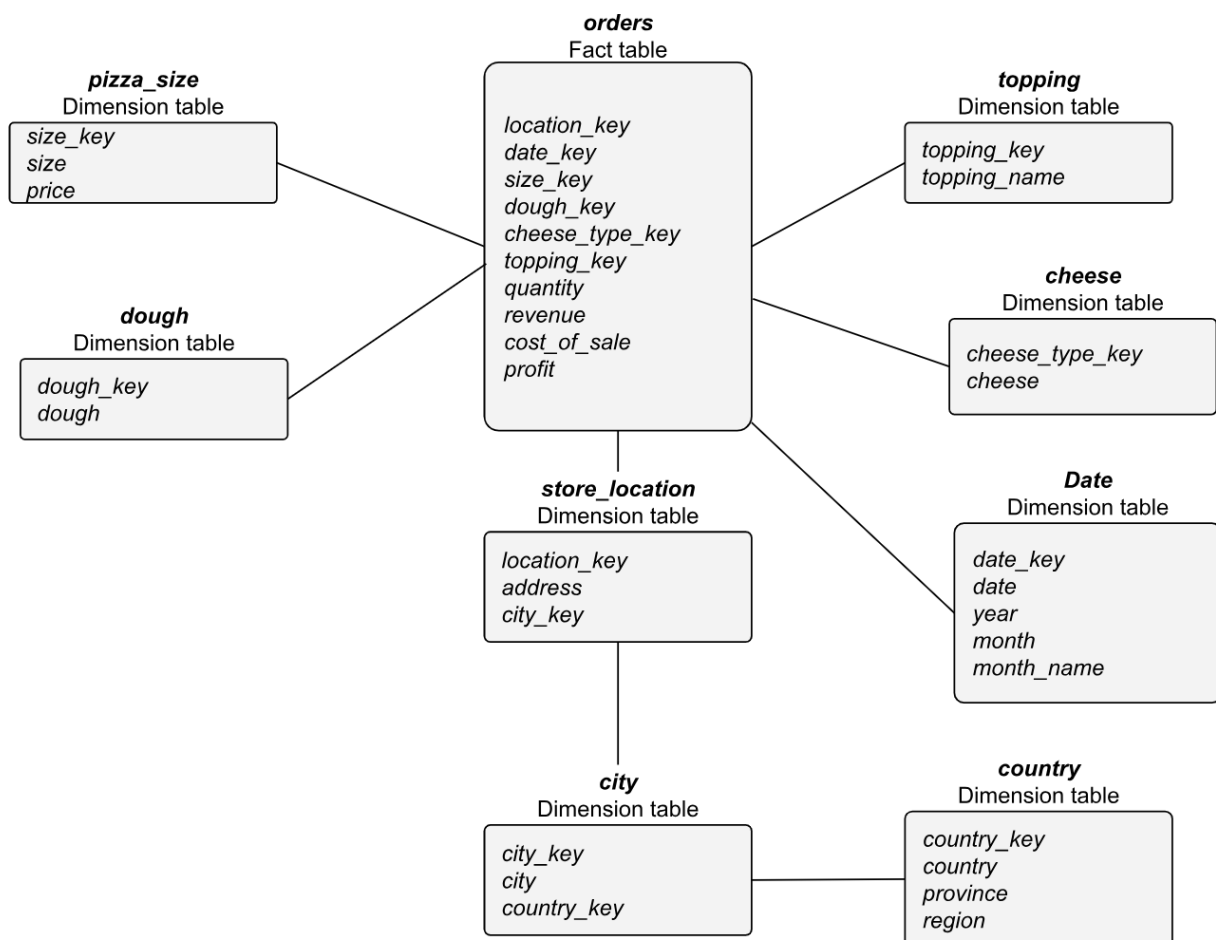
We included 'city' and 'country' as separate dimension tables because there are thousands of cities and hundreds of countries. Furthermore, the presence of region and province attributes would require the table to be further normalized and split.

The reason for linking the country dimension table to the city dimension table is to optimize the information retrieval from the data warehouse.

Again, note that we have assumed that the price of the pizza is dependent on only the size of the pizza ordered. Same is reflected in the dimension tables.

### Points to note:

- In the above diagram, we altered the fact table to add a few more attributes, which are calculated fields (i.e., revenue, profit and cost of sale), in order to improve the quality of the fact table and the ultimate analysis that would be carried out using it. We also used the same schema in our R files to generate our dimensions and facts.
- In addition, we introduce the price and cost\_of\_sale into the orders fact table, to calculate revenue and profit, respectively.



- We also introduced quantity and profit fields in order to provide deeper analysis of the orders. These fields are derived fields and are populated during fact population.
- Revenue is a value derived using price x quantity (in orders table).
- Profit is a value derived using revenue - cost of sale.

The R file attached provides the code that was used to build the OLAP.

The following sections provide the details of the tasks completed.

**Note:** Even though the keys for each dimension table are supposed to be numeric, for the purpose of better understanding our cube, we have replaced numeric values with strings.

### 1c. Generating data

We first created a date dimension having dates in the first 15 days of October 2023. This range can be expanded to generate dates of multiple months or years, and the code attached would simply split it into relevant months, years, month names and dates accordingly. We experimented with this and generated a sample dataset for 1 year of data, however, we proceeded with 15 days of data to keep our analysis straightforward and to increase the readability of our results.

Next, we populated the rest of our dimensions.

```

21 pizza_size <-
22   data.frame(size_key=c("P", "S", "M", "L", "XL"),
23             size=c("personal", "small", "medium", "large", "xlarge"),
24             price=c(5,15,20,25,30))
25
26 topping <-
27   data.frame(topping_key=c("tomatoes", "pepper", "onions", "pepperoni"),
28             topping_name=c("tomatoes", "pepper", "onions", "pepperoni"))
29
30 cheese <-
31   data.frame(cheese_key=c("swiss", "cheddar", "mozzarella", "parmesan"),
32             cheese=c("swiss", "cheddar", "mozzarella", "parmesan"))
33
34 dough <-
35   data.frame(dough_key=c("whole weat thin", "white regular", "stuffed crust", "regular"),
36             dough=c("whole weat thin", "white regular", "stuffed crust", "regular"))
37
38 store_location <-
39   data.frame(location_key=c("1", "2", "3", "4"),
40             address=c("336 Rideau St", "458 Rideau St", "471 Yonge St", "124 Fulton St"),
41             city_key=c("2", "2", "2", "1"))
42
43 city <-
44   data.frame(city_key = c("1", "2"),
45             city = c("New York", "Ottawa"),
46             country_key = c("1", "2"))
47
48 country <-
49   data.frame(country_key = c("1", "2"),
50             country = c("United States of America", "Canada"),
51             province = c("New York", "Ontario"),
52             region = c("region a", "region b"))
53
54
55 profit <-
56   data.frame(profit_key=c("1", "2", "3"),
57             profit=c(22.5, 5.70, 11.20),
58             revenue = c(45, 12, 32))
59
60 quantity <- c(1,2,3)
61

```

All of the dimensions were with respective fields. In the case of topping, cheese, and dough our primary key was the same as the name to enhance the readability of our cube. In the case where we added the numeric values to the key, we proceeded to merge the fact with the respective dimension to retrieve the non numeric values (cheese names, dough name etc.) to improve readability. However, we have not included this into our code at this point.

The quantity field was introduced here but used while sampling.

Next step was to create a function to generate orders which would populate our Orders fact. We completed this by using the following code:

```

66 # Function to generate the orders table
67 gen_orders <- function(no_of_rec) {
68   # Generate transaction data randomly
69   location_ <- sample(store_location$location_key, no_of_rec, replace=T, prob=c(2,1,1,2))
70   order_date_ <- sample(date$date_key, no_of_rec, replace=T)
71   order_month_ <- sample(date$Month_Name, no_of_rec, replace=T)
72   size_ <- sample(pizza_size$size_key, no_of_rec, replace=T, prob=c(1, 3, 2, 3, 1))
73   dough_ <- sample(dough$dough_key, no_of_rec, replace=T, prob=c(1, 3, 1, 3))
74   cheese_ <- sample(cheese$cheese_key, no_of_rec, replace=T, prob=c(2, 3, 1, 2))
75   topping_ <- sample(topping$topping_key, no_of_rec, replace=T, prob=c(3, 2, 1, 1))
76   #profit_ <- sample(profit$profit_key, no_of_rec, replace=T, prob=c(1, 3, 2))
77   quantity_ <- sample(quantity, no_of_rec, replace=T, prob=c(3, 2, 1))
78   price <- pizza_size$price[match(size_, pizza_size$size_key)]
79   revenue <- price * quantity_
80   cost_of_sales <- (sample(c(0.8,0.85), no_of_rec, replace=T, prob=c(3, 2))) * revenue
81   profit_ <- revenue - cost_of_sales
82
83
84   orders <- data.frame(location_      # = location_,
85                        ,order_date_  # = order_date_,
86                        ,order_month_
87                        ,size_        # = size_,
88                        ,dough_       # = dough_,
89                        ,cheese_      # = cheese_,
90                        ,topping_     # = topping_,
91                        #,profit_     # = profit_,
92                        ,quantity_    # = quantity
93                        ,price
94                        ,revenue
95                        ,cost_of_sales
96                        ,profit_
97                        )
98
99   # Sort the records by time order
100  orders <- orders[order(orders$order_date_),]
101  row
102  return(orders)
103 }
104
105 # Creating the orders_fact using function
106 orders_fact <- gen_orders(500)
107

```

The fields revenue, cost\_of\_sales and profit\_ are all calculated fields which have been calculated while populating the fact. This would certainly help in keeping the formulas for all of these fields constant. Even though this might be a costly activity in terms of resources, this approach makes the most sense since we are generating our own data and do not have data from a database providing us with these values.

2. Once generated, we used the Orders fact to build our revenue cube.

This revenue\_cube had the dimensions of size, order date, location, dough, cheese, topping and order month. To analyze other aspects of the order, for example profits, quantities sold or cost\_of\_sales, other cubes can also be built. We have included some of these in our code which we used for verification of our cube.

```

114 # Build up a cube for revenue
115 revenue_cube <-
116   tapply(orders_fact$revenue,
117         orders_fact[,c("size_", "order_date_", "location_", "dough_", "cheese_", "topping_", "order_month_")]
118         function(x){return(sum(x))})
119
120 # Showing the cells of the cube
121 revenue_cube
122

```

With our OLAP cube complete, we moved on to analyzing the data in it.

We started with applying a series of roll-up operations to bifurcate the fact data into dimension, leading us to finding some extremely interesting trends in our dataset.

```

159 #3.
160
161 #Roll-up 1 - Revenue in terms of size and topping
162 apply(revenue_cube, c("size_", "topping_"),
163       FUN=function(x) {return(sum(x, na.rm=TRUE))})
164
165 #Roll-up 2 - Revenue in terms of size and dough
166 apply(revenue_cube, c("size_", "dough_"),
167       FUN=function(x) {return(sum(x, na.rm=TRUE))})
168
169 #Roll-up 3 - Revenue in terms of size and cheese
170 apply(revenue_cube, c("size_", "cheese_"),
171       FUN=function(x) {return(sum(x, na.rm=TRUE))})
172
173
174 #Roll-up 4 - Revenue in terms of dough and cheese
175 apply(revenue_cube, c("dough_", "cheese_"),
176       FUN=function(x) {return(sum(x, na.rm=TRUE))})
177
178
179 #Roll-up 5 - Revenue in terms of dough and cheese
180 apply(revenue_cube, c("dough_", "topping_"),
181       FUN=function(x) {return(sum(x, na.rm=TRUE))})
182
183
184 #Roll-up 6 - Revenue in terms of size and location
185 apply(revenue_cube, c("size_", "location_"),
186       FUN=function(x) {return(sum(x, na.rm=TRUE))})
187

```

```

> #Roll-up 1 - Revenue in terms of size and topping
> apply(revenue_cube, c("size_", "topping_"),
+       FUN=function(x) {return(sum(x, na.rm=TRUE))})
      topping_
size_ onions pepper pepperoni tomatoes
L      1050    1600      1025    2050
M       540     920       560    1480
P       115     105        50     185
S       570     915       780    1665
XL      420     780       450    1260
>
> #Roll-up 2 - Revenue in terms of size and dough
> apply(revenue_cube, c("size_", "dough_"),
+       FUN=function(x) {return(sum(x, na.rm=TRUE))})
      dough_
size_ regular stuffed crust white regular whole weat thin
L      1825      800      2175      925
M      1260      400      1360      480
P       215       30       165       45
S      1260      510      1635      525
XL     1350      330       900      330
>
> #Roll-up 3 - Revenue in terms of size and cheese
> apply(revenue_cube, c("size_", "cheese_"),
+       FUN=function(x) {return(sum(x, na.rm=TRUE))})
      cheese_
size_ cheddar mozzarella parmesan swiss
L      2225      850      1525    1125
M      1520      460       560    960
P       150       30       150    125
S      1260      555      1140    975
XL     1380      570       420    540
>

```

```

>
> #Roll-up 4 - Revenue in terms of dough and cheese
> apply(revenue_cube, c("dough_", "cheese_"),
+       FUN=function(x) {return(sum(x, na.rm=TRUE))})
      cheese_
dough_    cheddar mozzarella parmesan swiss
regular      2350      870     1245  1445
stuffed crust   815      275      540   440
white regular  2280     1010     1625  1320
whole weat thin 1090      310      385   520
>
>
> #Roll-up 5 - Revenue in terms of dough and cheese
> apply(revenue_cube, c("dough_", "topping_"),
+       FUN=function(x) {return(sum(x, na.rm=TRUE))})
      topping_
dough_    onions pepper pepperoni tomatoes
regular      1095     1530      860   2425
stuffed crust   510      480      345    735
white regular   895     1675     1160   2505
whole weat thin  195      635      500    975
>
>
> #Roll-up 6 - Revenue in terms of size and location
> apply(revenue_cube, c("size_", "location_"),
+       FUN=function(x) {return(sum(x, na.rm=TRUE))})
      location_
size_    1    2    3    4
L    2125 1075 800 1725
M    1000  380 800 1320
P     160   75  55  165
S     915  615 885 1515
XL    660  660 540 1050
>

```

With our roll-up operations, we quickly analyzed the low revenue generated by “Personal” pizza size across all dimensions. And in every roll-up operation, the highest revenue was brought in by “Large” pizza size orders.

### Roll-up analysis

1. In the size vs toppings roll-up, Large pizzas with tomatoes generated the highest revenue, while Personal pizzas with pepperoni topping was the least popular. Overall, the tomato toppings generated the most revenue across all pizza sizes.
2. Similarly, in size vs dough comparison, we noticed that regular and white regular dough performed equally as well, competing with each other very closely. While both stuffed crust and whole weat thin lagged behind, with stuffed crust performing the worst.
3. The most interesting roll-up results were brought by rollup 6 comparing size and location. It seemed that while all locations performed almost equally well, the Personal pizza performed the worst in all locations.
4. It is worth noting, that in all of the roll up operations performed, personal pizzas were the worst performing pizza sizes. Where as the large and extra large pizzas were the most sold regardless of the location, topping, order data. Therefore, customers are preferring larger pizzas over small or medium sized pizzas.

```

#Drill-down 1 - Revenue in terms of order_month_ and location_ and size_
apply(revenue_cube, c("order_month_", "location_", "size_"),
      FUN=function(x) {return(sum(x, na.rm=TRUE))})

#Drill-down 2 - Revenue in terms of order_month_ and topping_ and size_
apply(revenue_cube, c("order_month_", "topping_", "size_"),
      FUN=function(x) {return(sum(x, na.rm=TRUE))})

#Drill-down 3 - Revenue in terms of order_month_ and topping_ and size_
apply(revenue_cube, c("order_month_", "location_", "size_"),
      FUN=function(x) {return(sum(x, na.rm=TRUE))})

#Drill-down 4 - Revenue in terms of order_month_ and topping_ and size_
apply(revenue_cube, c("dough_", "cheese_", "size_"),
      FUN=function(x) {return(sum(x, na.rm=TRUE))})

#Drill-down 5 - Revenue in terms of order_month_ and topping_ and size_
apply(revenue_cube, c("size_", "dough_", "location_"),
      FUN=function(x) {return(sum(x, na.rm=TRUE))})

> #Drill-down 1 - Revenue in terms of order_month_ and location_ and size_
> apply(revenue_cube, c("order_month_", "location_", "size_"),
+       FUN=function(x) {return(sum(x, na.rm=TRUE))})
, , size_ = L

      location_
order_month_  1  2  3  4
      October 2125 1075 800 1725

, , size_ = M

      location_
order_month_  1  2  3  4
      October 1000 380 800 1320

, , size_ = P

      location_
order_month_  1  2  3  4
      October 160 75 55 165

, , size_ = S

      location_
order_month_  1  2  3  4
      October 915 615 885 1515

, , size_ = XL

      location_
order_month_  1  2  3  4
      October 660 660 540 1050

```

```

> #Drill-down 2 - Revenue in terms of order_month_ and topping_ and size_
> apply(revenue_cube, c("order_month_", "topping_", "size_"),
+       FUN=function(x) {return(sum(x, na.rm=TRUE))})
, , size_ = L
      topping_
order_month_ onions pepper pepperoni tomatoes
      october   1050   1600      1025      2050
, , size_ = M
      topping_
order_month_ onions pepper pepperoni tomatoes
      october    540    920      560      1480
, , size_ = P
      topping_
order_month_ onions pepper pepperoni tomatoes
      october    115    105       50       185
, , size_ = S
      topping_
order_month_ onions pepper pepperoni tomatoes
      october    570    915      780      1665
, , size_ = XL
      topping_
order_month_ onions pepper pepperoni tomatoes
      october    420    780      450      1260

> #Drill-down 3 - Revenue in terms of order_month_ and topping_ and size_
> apply(revenue_cube, c("order_month_", "location_", "size_"),
+       FUN=function(x) {return(sum(x, na.rm=TRUE))})
, , size_ = L
      location_
order_month_  1  2  3  4
      october 2125 1075 800 1725
, , size_ = M
      location_
order_month_  1  2  3  4
      october 1000 380 800 1320
, , size_ = P
      location_
order_month_  1  2  3  4
      october 160 75 55 165
, , size_ = S
      location_
order_month_  1  2  3  4
      october 915 615 885 1515
, , size_ = XL
      location_
order_month_  1  2  3  4
      october 660 660 540 1050

```



```
> #Drill-down 4 - Revenue in terms of order_month_ and topping_ and size_
> apply(revenue_cube, c("dough_", "cheese_", "size_"),
+       FUN=function(x) {return(sum(x, na.rm=TRUE))})
, , size_ = L

      dough_      cheese_
      cheddar mozzarella parmesan swiss
regular      825      275      400      325
stuffed crust 275        25      325      175
white regular 675      350      650      500
whole weat thin 450      200      150      125

, , size_ = M

      dough_      cheese_
      cheddar mozzarella parmesan swiss
regular      460      160      200      440
stuffed crust 200        40      100      60
white regular 640      240      220      260
whole weat thin 220        20        40      200

, , size_ = P

      dough_      cheese_
      cheddar mozzarella parmesan swiss
regular        75        15        75      50
stuffed crust   10         0        10      10
white regular   50        15        35      65
whole weat thin  15         0        30        0

, , size_ = S

      dough_      cheese_
      cheddar mozzarella parmesan swiss
regular      390        90      360      420
stuffed crust 180      180      105      45
white regular 525      195      570      345
whole weat thin 165        90      105      165

, , size_ = XL

      dough_      cheese_
      cheddar mozzarella parmesan swiss
regular      600      330      210      210
stuffed crust 150        30         0      150
white regular 390      210      150      150
whole weat thin 240         0        60        30
```

---

```
> #Drill-down 5 - Revenue in terms of order_month_ and topping_ and size_
> apply(revenue_cube, c("size_", "dough_", "location_"),
+       FUN=function(x) {return(sum(x, na.rm=TRUE))})
, , location_ = 1

      dough_
size_ regular stuffed crust white regular whole weat thin
L      800      275      850      200
M      380        60      360      200
P       70         20        60       10
S      180      135      495      105
XL     240      150      120      150

, , location_ = 2

      dough_
size_ regular stuffed crust white regular whole weat thin
L      475      100      300      200
M      180        60        80       60
P       30         0        35       10
S      150      150      225       90
XL     360        90      210        0

, , location_ = 3

      dough_
size_ regular stuffed crust white regular whole weat thin
L      150      125      325      200
M      300      100      400         0
P       35         0        20         0
S      285        75      345      180
XL     180        30      270        60

, , location_ = 4

      dough_
size_ regular stuffed crust white regular whole weat thin
L      400      300      700      325
M      400      180      520      220
P       80       10        50       25
S      645      150      570      150
XL     570        60      300      120
```

## Drill-down analysis

1. Through our drill down operations, we were able to find staggering analysis patterns. In our first operation we drilled down into the OLAP cube on the 3 dimensions of order month, location and size. We found that as per our roll-up operations, Personal pizzas were sold extremely less than any other pizza. Therefore, the store should not focus on getting more components of Personal sized pizza and instead buy large and extra large in bulk quantities as they prove to be the most selling and revenue generating sizes.
2. When we drilled down into order month vs topping vs size, we analyzed that pepperoni and onion generated the least revenue in the entire month for all pizza sizes. Whereas tomatoes were the most favored. It would thus be advised to the store to buy more tomato topping components than pepperoni. Pepper toppings were closely following tomatoes and would require to be replenished.
3. Our most important finding was in drilling down dough vs cheese vs size. In this operation we noticed that pizzas with mozzarella cheese and stuffed crust were never ordered. Same was the case with whole wheat thin dough and mozzarella and swiss cheese. Therefore, these components should not be bought or bought in very less amounts as they did not generate ANY revenue for the store. Furthermore regular dough with cheddar cheese is the most selling pizza, and must be replenished in high amounts as that would generate the highest revenue for the store.
4. Much like our roll0up analysis, the drop-down operations confirmed that the personal pizza sizes were the least bought. Whereas the large and extra large were generating the most revenue. It can thus be concluded that the large and XL pizzas are preferred by the customers and the store should prefer buying components for these sizes.

**Note:** Please note that, in Part A, question 1C and 2 (first part), the data was generated using R, as per the confirmation on email given below. We are mentioning this point in this report just to make sure there is no confusion in the assignment question and the corresponding responses presented in this document.

**Re: DTI 5126 Assignment 2 questions**  
mailto:lpaiv023@uottawa.ca

Pouya Khodaei <pkhod015@uottawa.ca>  
Sat 11/4/2023 7:40 PM  
To: Lakshika Josephine Paiva <lpaiv023@uottawa.ca>  
Cc: Usman Bashir <mbash028@uottawa.ca>

Hi,

Yes, just get the idea from tutorial material and do not copy-paste. No .CSV file is available for the first part.

Best,  
Pouya

---

**From:** Lakshika Josephine Paiva <lpaiv023@uottawa.ca>  
**Date:** Saturday, November 4, 2023 at 6:06 PM  
**To:** Pouya Khodaei <pkhod015@uottawa.ca>  
**Cc:** Usman Bashir <mbash028@uottawa.ca>  
**Subject:** Re: DTI 5126 Assignment 2 questions

Hi Pouya,

Thanks.  
For Part A, Q1-c, can we generate the data for the dimension files using the data generation code taught in tutorial 2?

Or should we just create it manually/ offline in CSV files and then read them using R code?

Thanks and regards,  
Lakshika

## Part B

Prior to the tasks in the assignment, the data was imported into RStudio. Then the .csv file with delimiters, was converted to a table to ensure better readability in rows and columns. R code used for the process is shown:

Import:

```
> library(readr)
> bank_additional_full <- read_csv("bank-additional-full.csv")
Rows: 41188 Columns: 1
— Column specification —————
Delimiter: ", "
chr (1): age;job;marital;education;default;housing;loan;contact;month;d...
```

Imported data:

The screenshot shows the RStudio interface with a file named 'bank\_additional\_full.csv' open. The data is displayed as a single column of text, where each row contains a semicolon-separated list of values for all 21 columns. The first few rows are visible:

Row	Value
1	56;housemaid;married;basic.4y;no;no;no;te...
2	57;services;married;high.school;unknown;no;...
3	37;services;married;high.school;no;yes;no;t...
4	40;admin.;married;basic.6y;no;no;no;teleph...
5	56;services;married;high.school;no;no;yes;t...
6	45;services;married;basic.9y;unknown;no;no;...
7	59;admin.;married;professional.course;no;no;...
8	41;blue-collar;married;unknown;unknown;no;...

Showing 1 to 9 of 41,188 entries, 1 total columns

Converted data into a table and assigned to a data frame.

```
> bank_additional.df = read.csv(file="bank-additional-full.csv", sep=";", header=T)
```

Converted data:

The screenshot shows the RStudio interface with the data converted into a table with 21 columns. The first few rows are visible:

	age	job	marital	education	default	housing	loan	contact	month	day_of_w
1	56	housemaid	married	basic.4y	no	no	no	telephone	may	mon
2	57	services	married	high.school	unknown	no	no	telephone	may	mon
3	37	services	married	high.school	no	yes	no	telephone	may	mon
4	40	admin.	married	basic.6y	no	no	no	telephone	may	mon
5	56	services	married	high.school	no	no	yes	telephone	may	mon
6	45	services	married	basic.9y	unknown	no	no	telephone	may	mon
7	59	admin.	married	professional.course	no	no	no	telephone	may	mon
8	41	blue-collar	married	unknown	unknown	no	no	telephone	may	mon

Showing 1 to 9 of 41,188 entries, 21 total columns

Answers to questions:

1. The data columns which are required were preserved and other columns were dropped.

```
> bank_simple.df <- bank_additional_full.df[-c(2:3, 5:12, 15:20)]
```

```
> View(bank_simple.df)
```

The screenshot shows the RStudio interface with the simplified data frame 'bank\_simple.df' open. The data is displayed as a table with 5 columns: age, education, pdays, previous, and y. The first 10 rows are visible:

	age	education	pdays	previous	y
1	56	basic.4y	999	0	no
2	57	high.school	999	0	no
3	37	high.school	999	0	no
4	40	basic.6y	999	0	no
5	56	high.school	999	0	no
6	45	basic.9y	999	0	no
7	59	professional.course	999	0	no
8	41	unknown	999	0	no
9	24	professional.course	999	0	no
10	25	high.school	999	0	no

Showing 1 to 11 of 41,188 entries, 5 total columns

In addition, the “y” column was renamed to “response” as it was not descriptive of the data in the column.

```
> names(bank_simple.df)[names(bank_simple.df) == "y"] <- "response"
```

	age	education	pdays	previous	response
1	56	basic.4y	999	0	no
2	57	high.school	999	0	no
3	37	high.school	999	0	no
4	40	basic.6y	999	0	no
5	56	high.school	999	0	no
6	45	basic.9y	999	0	no
7	59	professional.course	999	0	no
8	41	unknown	999	0	no

Showing 1 to 8 of 41,188 entries, 5 total columns

- The value “999” is replaced with “NA”. The ‘before’ and ‘after’ output is shown below.

Before:

```
> table(bank_simple.df$pdays)

 0    1   10   11   12   13   14   15   16   17   18   19 
15   26   52   28   58   36   20   24   11    8    7    3 
 2   20   21   22   25   26   27    3    4    5    6    7 
61    1    2    3    1    1    1  439  118   46  412   60 
 8    9   999 
18   64 39673
```

```
bank_simple.df$pdays[bank_simple.df$pdays == "999"] <- "NA"
```

After:

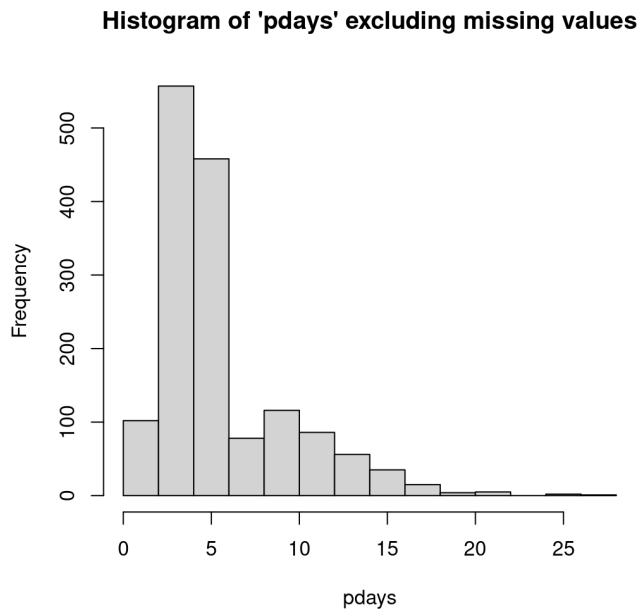
```
> table(bank_simple.df$pdays)

 0    1   10   11   12   13   14   15   16   17   18   19 
15   26   52   28   58   36   20   24   11    8    7    3 
 2   20   21   22   25   26   27    3    4    5    6    7 
61    1    2    3    1    1    1  439  118   46  412   60 
 8    9    NA 
18   64 39673
```

- In reality ‘999’ denotes missing values. Because the ‘pdays’ column is numeric, having 999 in comparison to the other values which are drastically smaller, will make ‘999’ look like part of the data and affect/skew the data in calculations. For example, if the mean is calculated including ‘999’ values, the result will take an abnormal value with 999 in the calculation, whereas actually it is significantly lesser than that. Hence, it is better to substitute it with NA.
- See the code and histogram below:

```
> bank_simple.df$pdays[bank_simple.df$pdays == 999] <- "NA"
> bank_simple.df$pdays <- as.numeric(as.character(bank_simple.df$pdays))
> hist(bank_simple.df$pdays[!is.na(bank_simple.df$pdays)], xlab="pdays",
ylab="Frequency", main="Histogram of pdays excluding missing values")
```

The histogram is self-explanatory, and it shows the frequency of pdays.



5. Converting categorical values to numerical values in the education column.

```
> bank_simple.df$education[bank_simple.df$education=="illiterate"]<-"0"
> bank_simple.df$education[bank_simple.df$education=="basic.4y"]<-"4"
> bank_simple.df$education[bank_simple.df$education=="basic.6y"]<-"6"
> bank_simple.df$education[bank_simple.df$education=="basic.9y"]<-"9"
> bank_simple.df$education[bank_simple.df$education=="high.school"]<-"12"
> bank_simple.df$education[bank_simple.df$education=="professional.course"]<-"12a"
> bank_simple.df$education[bank_simple.df$education=="university.degree"]<-"16"
> bank_simple.df$education[bank_simple.df$education=="unknown"]<-"NA"
> count(bank_simple.df$education)
```

	x	freq
1	0	18
2	12	9515
3	12a	5243
4	16	12168
5	4	4176
6	6	2292
7	9	6045
8	NA	1731

6. Following are the mean, median and mode calculations of the 'age' field:

```
> mean(bank_simple.df$age)
[1] 40.02406
> median(bank_simple.df$age)
[1] 38
> library(DescTools)
> Mode(bank_simple.df$age)
[1] 31
attr(,"freq")
[1] 1947
```

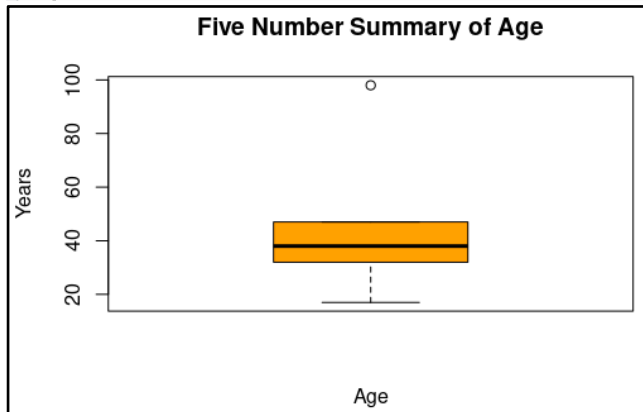
Please note that, 'mode' was calculated after using the library (DescTools)

Calculation of five number summary to show the data, and assigning it to a variable.

```
> age_fivenum <- fivenum(bank_simple.df$age)
> age_fivenum
[1] 17 32 38 47 98
```

Boxplot of five number summary of age:

```
> boxplot(age_fivenum,
+         data=age_fivenum,
+         main="Five Number Summary of Age",
+         xlab="Age",
+         ylab="Years",
+         col="orange",
+         border="black"
+ )
```

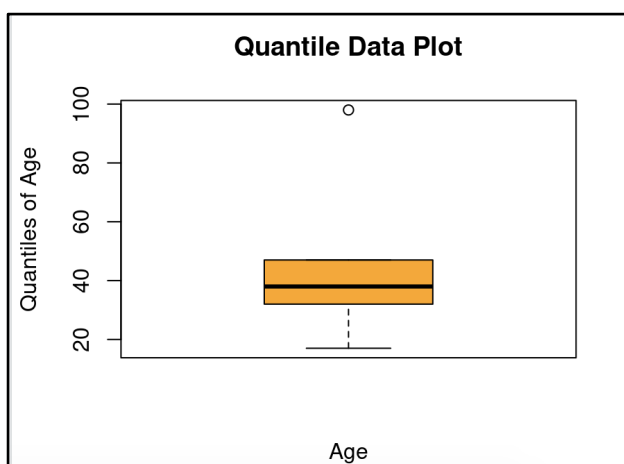


Quantile information:

```
> quantile(bank_simple.df$age)
 0%  25%  50%  75% 100%
17   32   38   47   98
```

Following is the boxplot for quantiles in age.

```
> boxplot(quantile_plot, data=quantile_plot, main="Quantile Data Plot",
+         xlab="Age", ylab="Quantiles of Age", col="orange", border="black")
```



7. Age variable standardized and saved in “age\_z”

```
> age_z <- scale(bank_simple.df$age)
```

## 8. Detecting outliers in age\_z.

```
> age_z <- scale(bank_simple.df$age)
> out <- boxplot.stats(age_z)$out
> out_ind <- which(age_z %in% c(out))
> bank_simple.df[out_ind, ]
```

According to the output, individuals aged 70 and above are considered outliers in this dataset.

	age	education	pdays	previous	response
27714	70	basic.4y	NA	0	yes
27758	76	university.degree	NA	0	no
27781	73	university.degree	NA	1	no
27801	88	basic.4y	NA	0	no
27803	88	basic.4y	NA	0	yes
27806	88	basic.4y	NA	0	yes
27809	88	basic.4y	NA	0	no
27811	88	basic.4y	NA	0	yes
27812	88	basic.4y	NA	0	yes
27813	88	basic.4y	NA	0	no
27814	88	basic.4y	NA	0	yes
27815	88	basic.4y	NA	0	no
27816	88	basic.4y	NA	0	no
27817	88	basic.4y	NA	0	yes
27818	88	basic.4y	NA	0	yes
27819	88	basic.4y	NA	0	yes
27827	95	basic.6y	NA	0	no
27838	70	basic.4y	NA	1	no
27839	70	basic.4y	NA	1	no
27845	70	basic.4y	NA	0	no
27852	77	unknown	NA	0	yes
27876	75	basic.9y	NA	0	no
27880	70	university.degree	NA	0	no
27903	70	basic.4y	NA	1	no
27931	73	university.degree	NA	0	yes
27951	80	basic.4y	NA	0	no
27952	80	basic.4y	NA	0	no
27964	80	professional.course	NA	0	yes
28221	72	basic.4y	NA	0	no
28222	72	basic.4y	NA	0	no
28313	82	unknown	NA	0	yes
28457	73	basic.4y	NA	0	yes
28505	71	basic.4y	NA	0	no
28531	70	basic.4y	NA	0	yes
28541	70	basic.4y	NA	1	yes
28587	70	basic.4y	NA	0	no
28620	71	high.school	NA	0	no
28733	70	unknown	NA	0	yes
28774	70	unknown	NA	0	no
29226	71	unknown	NA	0	yes
29264	75	basic.4y	NA	0	no
29499	73	basic.4y	6	1	no
29626	73	basic.4y	NA	0	no
29669	71	university.degree	NA	0	yes
29683	75	basic.4y	NA	0	yes
29974	75	basic.4y	NA	0	no
29978	78	basic.4y	NA	0	yes
29982	75	basic.4y	NA	1	yes
29988	70	basic.6y	NA	0	no

29991	78	basic.4y	NA	0	no
30001	75	basic.4y	NA	1	yes
30005	78	basic.4y	NA	0	yes
30007	85	basic.4y	NA	0	yes
30073	85	basic.4y	NA	0	no
30079	85	basic.4y	NA	0	no
30080	80	high.school	NA	3	no
30089	71	university.degree	NA	0	no
30104	85	basic.4y	NA	0	no
30111	85	basic.4y	NA	0	no
30134	79	basic.9y	NA	0	yes
30172	77	basic.4y	NA	1	no
30215	83	basic.4y	NA	0	no
30226	81	professional.course	NA	0	no
30228	71	university.degree	5	1	no
30242	81	professional.course	NA	0	no
30335	73	university.degree	NA	1	no
30336	71	basic.4y	NA	1	no
30391	71	basic.4y	NA	0	no
30431	88	high.school	NA	0	no
30461	81	basic.9y	NA	1	no
30590	81	basic.4y	NA	0	yes
35834	81	unknown	NA	0	yes
35849	71	unknown	NA	1	no
35857	83	basic.9y	3	3	no
35879	75	basic.4y	NA	0	no
35974	78	high.school	3	2	no
36184	88	basic.4y	NA	0	no
36286	77	basic.9y	NA	1	no
36312	72	university.degree	NA	0	no
36384	79	high.school	NA	1	no
36385	79	high.school	NA	1	no
36817	74	basic.4y	NA	0	yes
36999	75	basic.4y	NA	0	no
37137	72	university.degree	NA	0	no
37138	72	university.degree	NA	0	no
37171	70	basic.4y	NA	0	no
37187	79	unknown	NA	0	no
37191	74	high.school	NA	0	no
37193	74	high.school	NA	0	no
37194	74	high.school	NA	1	yes
37196	74	high.school	NA	0	no
37207	76	basic.4y	NA	0	yes
37208	76	basic.4y	NA	0	yes
37214	82	university.degree	NA	2	no
37220	75	basic.4y	NA	0	yes
37228	70	basic.4y	NA	0	yes
37236	73	basic.4y	NA	0	yes
37238	73	basic.9y	15	1	no
37240	73	basic.4y	NA	0	no
37258	73	basic.4y	NA	0	no
37261	76	university.degree	NA	0	yes
37317	70	basic.4y	NA	1	no
37342	85	professional.course	NA	0	yes
37356	80	illiterate	6	1	yes
37372	70	high.school	NA	0	no
37404	74	university.degree	NA	1	yes
37455	74	professional.course	NA	0	no
37456	76	basic.4y	NA	0	no
37473	88	basic.4y	NA	0	no
37480	74	basic.4y	NA	1	no



37494	81	basic.6y	4	1	no
37506	76	basic.6y	NA	1	no
37510	74	professional.course	NA	1	yes
37513	76	university.degree	NA	1	no
37526	73	professional.course	NA	0	no
37533	72	basic.4y	NA	1	no
37546	70	professional.course	NA	1	no
37569	71	basic.4y	NA	0	yes
37571	70	professional.course	NA	0	yes
37587	70	professional.course	NA	0	no
37598	76	professional.course	15	1	yes
37602	72	basic.4y	NA	0	no
37603	73	basic.4y	NA	0	yes
37605	80	high.school	NA	0	yes
37636	74	basic.9y	NA	1	yes
37662	71	basic.4y	NA	0	yes
37676	74	basic.4y	13	1	yes
37680	80	basic.4y	NA	0	no
37691	74	university.degree	NA	0	yes
37693	73	basic.4y	NA	0	no
37716	74	basic.9y	NA	1	no
37717	71	university.degree	6	2	yes
37736	76	basic.4y	NA	0	no
37737	76	basic.4y	NA	1	no
37744	87	basic.4y	NA	0	yes
37757	79	basic.4y	3	2	yes
37766	70	university.degree	NA	1	no
37770	74	basic.4y	6	1	yes
37776	88	university.degree	NA	0	no
37785	81	high.school	NA	0	no
37819	80	basic.4y	NA	0	yes
37820	80	basic.4y	3	2	no
37821	78	basic.4y	NA	0	no
37826	71	university.degree	NA	0	no
37827	71	university.degree	NA	0	no
37862	72	unknown	NA	1	yes
37869	73	university.degree	NA	0	no
37871	73	basic.4y	NA	1	yes
37874	73	high.school	NA	0	no
37906	79	basic.9y	NA	1	yes
37921	72	unknown	NA	0	no
37936	71	basic.4y	NA	0	no
37947	83	unknown	NA	0	no
37952	76	unknown	NA	0	no
37953	76	unknown	NA	0	yes
37955	72	basic.4y	3	2	yes
37959	71	professional.course	NA	0	no
37998	71	basic.4y	NA	0	yes
38000	76	basic.4y	NA	0	no
38006	75	unknown	NA	1	yes
38008	71	basic.4y	NA	0	no
38020	78	unknown	NA	0	no
38021	78	unknown	NA	1	yes
38023	91	university.degree	NA	2	no
38033	91	university.degree	NA	0	no
38034	76	basic.4y	3	1	yes
38046	73	basic.4y	NA	0	no
38053	76	university.degree	NA	1	no
38055	73	basic.4y	NA	0	no
38061	71	unknown	NA	0	yes
38066	83	unknown	NA	0	no

38072	70	basic.4y	3	2	yes
38075	70	basic.4y	NA	2	yes
38082	70	basic.4y	3	1	yes
38089	70	basic.4y	3	1	no
38126	70	high.school	NA	0	no
38128	70	university.degree	NA	0	yes
38130	70	university.degree	NA	0	no
38137	81	basic.4y	NA	0	yes
38145	70	basic.4y	NA	1	no
38146	70	basic.4y	NA	0	no
38167	78	basic.9y	NA	0	no
38170	71	basic.4y	NA	2	no
38176	71	basic.4y	NA	1	yes
38179	75	basic.4y	NA	0	yes
38180	83	professional.course	4	1	yes
38184	71	basic.9y	NA	0	no
38185	82	university.degree	NA	0	yes
38192	82	university.degree	NA	1	no
38193	82	university.degree	NA	0	no
38194	80	basic.4y	NA	0	no
38196	80	basic.4y	NA	0	no
38207	86	basic.4y	NA	0	no
38230	77	unknown	NA	0	no
38242	75	basic.9y	NA	0	no
38247	77	university.degree	NA	0	yes
38248	70	university.degree	NA	1	no
38253	80	basic.4y	NA	1	no
38254	71	university.degree	NA	0	no
38255	71	university.degree	NA	0	no

[ reached 'max' / getOption("max.print") -- omitted 269 rows ]